

# Lecture 28 — Causal Profiling, Memory Profiling

Patrick Lam & Jeff Zarnett

`patrick.lam@uwaterloo.ca jzarnett@uwaterloo.ca`

Department of Electrical and Computer Engineering  
University of Waterloo

December 3, 2022

As with queueing theory:

$$\text{allocs} > \text{frees} \implies \text{usage} \rightarrow \infty$$

At least more paging, maybe total out-of-memory.

But! Memory isn't really lost: we could free it.

Our tool for this comes from the Valgrind tool suite.

DHAT tracks allocation points and what happens to them over a program's execution.

An allocation point is a point in the program which allocates memory.

```

AP 1.1.1.1/2 {
  Total:    31,460,928 bytes (2.32%, 1,565.9/Minstr) in 262,171 blocks (4.41%, 13.05/Minstr)
           avg size 120 bytes, avg lifetime 986,406,885.05 instrs (4.91% of program duration)
  Max:      16,779,136 bytes in 65,543 blocks, avg size 256 bytes
  At t-gmax: 0 bytes (0%) in 0 blocks (0%), avg size 0 bytes
  At t-end:  0 bytes (0%) in 0 blocks (0%), avg size 0 bytes
  Reads:     5,964,704 bytes (0.11%, 296.88/Minstr), 0.19/byte
  Writes:    10,487,200 bytes (0.51%, 521.98/Minstr), 0.33/byte
  Allocated at {
    ^1: 0x95CACC9: alloc (alloc.rs:72)
        [omitted]
    ^7: 0x95CACC9: parse_token_trees_until_close_delim (tokentrees.rs:27)
    ^8: 0x95CACC9: syntax::parse::lexer::tokentrees::<impl syntax::parse::lexer::
                StringReader<'a>>::parse_token_tree (tokentrees.rs:81)
    ^9: 0x95CAC39: parse_token_trees_until_close_delim (tokentrees.rs:26)
    ^10: 0x95CAC39: syntax::parse::lexer::tokentrees::<impl syntax::parse::lexer::
                StringReader<'a>>::parse_token_tree (tokentrees.rs:81)
    #11: 0x95CAC39: parse_token_trees_until_close_delim (tokentrees.rs:26)
    #12: 0x95CAC39: syntax::parse::lexer::tokentrees::<impl syntax::parse::lexer::
                StringReader<'a>>::parse_token_tree (tokentrees.rs:81)
  }
}

```

Going up the tree, DHAT tells you about all of the allocation points that share a call-stack prefix.

The ^ before the number indicates that the line is copied from the parent.

A # indicates that the line is unique to that node.

A sibling of the example above would diverge at the call on line 10.

# Shieldmaiden to Thor



What does Massif do?

- How much heap memory is your program using?
- How did this happen?

Next up: example from Massif docs.

# Example Allocation Program

---

```
fn g() {  
    let a = Vec::<u8>::with_capacity(4000);  
    std::mem::forget(a)  
}  
  
fn f() {  
    let a = Vec::<u8>::with_capacity(2000);  
    std::mem::forget(a);  
    g()  
}  
  
fn main() {  
  
    let mut a = Vec::with_capacity(10);  
    for _i in 0..10 {  
        a.push(Box::new([0;1000]))  
    }  
    f();  
    g();  
}
```

---



After we compile, run the command:

```
plam@amqui ~/c/p/l/l/L/alloc> valgrind --tool=massif target/debug/alloc
==406569== Massif, a heap profiler
==406569== Copyright (C) 2003-2017, and GNU GPL'd, by Nicholas Nethercote
==406569== Using Valgrind-3.16.1 and LibVEX; rerun with -h for copyright info
==406569== Command: target/debug/alloc
==406569==
==406569==
```

What happened?

- 1 The program ran slowly (because Valgrind!)
- 2 No summary data on the console  
(like memcheck or helgrind or cachegrind.)


Weird. What we got instead was the file `massif.out.[PID]`.

```
massif.out.[PID]:  
    plain text, sort of readable.
```

Better: `ms_print`.

Which has nothing whatsoever to do with Microsoft.  
Promise.

# Post-Processed Output



For a long time, nothing happens, then...kaboom!

Why? We gave it a trivial program.

We should tell Massif to care more about bytes than CPU cycles,  
with `--time-unit=B`.

Let's try that.

Run valgrind (Memcheck) first and make it happy before we go into figuring out where heap blocks are going with Massif.

Okay, what to do with the information from Massif, anyway?

Easy!

- Start with peak (worst case scenario) and see where that takes you (if anywhere).
- You can probably identify some cases where memory is hanging around unnecessarily.

Memory usage climbing over a long period of time, perhaps slowly, but never decreasing—memory filling with junk?

Large spikes in the graph—why so much allocation and deallocation in a short period?

- stack allocation ( - - stacks=yes).
- children of a process  
(anything split off with fork) if desired.
- low level stuff: if going beyond malloc, calloc, new, etc. and using mmap or brk that is usually missed, can do profiling at page level ( - - pages-as-heap=yes).



As is often the case,  
we have examined the tool on a trivial program.

Let's see if we can do some  
live demos of Massif at work.