

Lecture 30 — Clusters & Cloud Computing

Patrick Lam

`patrick.lam@uwaterloo.ca`

Department of Electrical and Computer Engineering
University of Waterloo

November 13, 2020


Part I

MPI & Clusters

So far, we've seen how to make things fast on one computer:

- threads;
- compiler optimizations;
- GPUs.

To get a lot of bandwidth, though, you need lots of computers,
(if you're lucky and the problem allows!)



Today: programming for performance with multiple computers via MPI.

Key Idea: Explicit Communication

Mostly we've seen shared-memory systems;
complication: must manage contention.

Recently, GPU programming: explicitly copy data.

Message-passing: yet another paradigm.

Message Passing Interface:

A language-independent communication protocol for parallel computers.

- Use it to run the same code on a number of **nodes** (different hardware threads; or servers in a cluster).
- Provides explicit message passing between nodes.
- Is the dominant model for high performance computing (de-facto standard).

MPI is a type of SPMD (single process, multiple data).

Idea: have multiple instances of the same program,
all working on different data.



MPI facilitates communication of data between processes.

```
// Initialize MPI
int MPI_Init(int *argc, char **argv)

// Determine number of processes within a communicator
int MPI_Comm_size(MPI_Comm comm, int *size)

// Determine processor rank within a communicator
int MPI_Comm_rank(MPI_Comm comm, int *rank)

// Exit MPI (must be called last by all processors)
int MPI_Finalize()

// Send a message
int MPI_Send (void *buf, int count, MPI_Datatype datatype,
               int dest, int tag, MPI_Comm comm)

// Receive a message
int MPI_Recv (void *buf, int count, MPI_Datatype datatype,
               int source, int tag, MPI_Comm comm,
               MPI_Status *status)
```

- `MPI_Comm`: a **communicator**,
often `MPI_COMM_WORLD` for global channel.
- `MPI_Datatype`: just an enum, e.g. `MPI_FLOAT_INT`, etc.
- `dest/source`: “rank” of the process (in a communicator)
to send a message to/receive a message from;
you may use `MPI_ANY_SOURCE` in `MPI_Recv`.
- Both `MPI_Send` and `MPI_Recv` are blocking calls—
see `man MPI_Send` or `man MPI_Recv` for more details.
- The `tag` allows you to organize your messages,
so you can filter all but a specific tag.

As with OpenCL kernels:

first, figure out what “current” process is supposed to compute.

```
// http://www.dartmouth.edu/~rc/classes/intro_mpi/  
#include <stdio.h>  
#include <mpi.h>  
  
int main (int argc, char * argv[])  
{  
    int rank, size;  
  
    /* start MPI */  
    MPI_Init (&argc, &argv);  
    /* get current process id */  
    MPI_Comm_rank (MPI_COMM_WORLD, &rank);  
    /* get number of processes */  
    MPI_Comm_size (MPI_COMM_WORLD, &size);  
    printf("Hello world from process %d of %d\n", rank, size);  
    MPI_Finalize();  
    return 0;  
}
```

Longer MPI Example (from Wikipedia)

Here's a common example:

- The “master” (rank 0) process creates some strings and sends them to the worker processes.
- The worker processes modify their string and send it back to the master.

Source:

http://en.wikipedia.org/wiki/Message_Passing_Interface.

```
/*  
  "Hello World" MPI Test Program  
*/  
#include <mpi.h>  
#include <stdio.h>  
#include <string.h>  
  
#define BUFSIZE 128  
#define TAG 0  
  
int main(int argc, char *argv[])  
{  
    char idstr[32];  
    char buff[BUFSIZE];  
    int numprocs;  
    int myid;  
    int i;  
    MPI_Status stat;
```

```
/* all MPI programs start with MPI_Init; all 'N'
 * processes exist thereafter
 */
MPI_Init(&argc,&argv);

/* find out how big the SPMD world is */
MPI_Comm_size(MPI_COMM_WORLD, &numprocs);

/* and this processes' rank is what? */
MPI_Comm_rank(MPI_COMM_WORLD, &myid);

/* At this point, all programs are running equivalently;
 * the rank distinguishes the roles of the programs in
 * the SPMD model, with rank 0 often used specially...
 */
```

```
if (myid == 0)
{
    printf("%d: We have %d processors\n", myid, numprocs);
    for (i=1; i<numprocs; i++)
    {
        sprintf(buff, "Hello %d! ", i);
        MPI_Send(buff, BUFSIZE, MPI_CHAR, i, TAG,
                 MPI_COMM_WORLD);
    }
    for (i=1; i<numprocs; i++)
    {
        MPI_Recv(buff, BUFSIZE, MPI_CHAR, i, TAG,
                 MPI_COMM_WORLD, &stat);
        printf("%d: %s\n", myid, buff);
    }
}
```

```
else
{
    /* receive from rank 0: */
    MPI_Recv(buff, BUFSIZE, MPI_CHAR, 0, TAG,
             MPI_COMM_WORLD, &stat);
    sprintf(idstr, "Processor %d ", myid);
    strncat(buff, idstr, BUFSIZE-1);
    strncat(buff, "reporting for duty", BUFSIZE-1);
    /* send to rank 0: */
    MPI_Send(buff, BUFSIZE, MPI_CHAR, 0, TAG,
             MPI_COMM_WORLD);
}

/* MPI Programs end with MPI Finalize; this is a weak
 * synchronization point.
 */
MPI_Finalize();
return 0;
}
```

```
// Wrappers for gcc (C/C++)  
mpicc  
mpicxx
```

```
// Compiler Flags  
OMPI_MPICC_CFLAGS  
OMPI_MPICXX_CXXFLAGS
```

```
// Linker Flags  
OMPI_MPICC_LDFLAGS  
OMPI_MPICXX_LDFLAGS
```

OpenMPI does not recommend that you set the flags yourself.
To see them, try:

```
# Show the flags necessary to compile MPI C applications  
shell$ mpicc —showme:compile
```

```
# Show the flags necessary to link MPI C applications  
shell$ mpicc —showme:link
```

```
mpirun -np <num_processors> <program>  
mpiexec -np <num_processors> <program> # a synonym
```

Starts `num_processors` instances of the program using MPI.

```
jon@riker examples master % mpicc hello_mpi.c  
jon@riker examples master % mpirun -np 8 a.out  
0: We have 8 processors  
0: Hello 1! Processor 1 reporting for duty  
0: Hello 2! Processor 2 reporting for duty  
0: Hello 3! Processor 3 reporting for duty  
0: Hello 4! Processor 4 reporting for duty  
0: Hello 5! Processor 5 reporting for duty  
0: Hello 6! Processor 6 reporting for duty  
0: Hello 7! Processor 7 reporting for duty
```

- By default, MPI uses the lowest-latency communication resource available; shared memory, in this case.



Matrix Multiplication



Neural Network

It be like that

MPI Matrix Multiplication Example

Highlights of: <http://www.nccs.gov/wp-content/training/mpi-examples/C/matmul.c>.

To compute the matrix product AB :

- 1 Initialize MPI.
- 2 If the current process is the master task (task id 0):
 - 1 Initialize matrices.
 - 2 Send work to each worker task:
row number (offset); number of rows; row contents from A ; complete contents of matrix B .

```
MPI_Send(&a[offset][0], rows*NCA, MPI_DOUBLE, dest,  
        mtype, MPI_COMM_WORLD);
```
 - 3 Wait for results from all worker tasks (`MPI_Recv`).
 - 4 Print results.
- 3 For all other tasks:
 - 1 Receive offset, number of rows, partial matrix A , and complete matrix B , using `MPI_Recv`:

```
MPI_Recv(&offset, 1, MPI_INT, MASTER,  
        mtype, MPI_COMM_WORLD, &status);
```
 - 2 Do the computation.
 - 3 Send the results back to the sender.

We can use nodes on a network (by using a hostfile).

We can even use MPMD:

- *multiple processes, multiple data*

```
% mpirun -np 2 a.out : -np 2 b.out
```

This launches a single parallel application.

- All in the same `MPI_COMM_WORLD`; but
- Ranks 0 and 1 are instances of `a.out`, and
- Ranks 2 and 3 are instances of `b.out`.

You could also use the `-app` flag with an appfile instead of typing out everything.

Your bottleneck for performance here is message-passing.


Keep the communication to a minimum!

In general, the more machines/farther apart they are, the slower the communication.

Each step from multicore machines to GPU programming to MPI triggers an order-of-magnitude decrease in communication bandwidth and similar increase in latency.

Part II

Cloud Computing



Historically:

- find \$\$\$;
- buy and maintain pile of expensive machines.

Not anymore!

We'll talk about Amazon's Elastic Compute Cloud (EC2)
and principles behind it.

You want a server on the Internet.

- Once upon a time: physical machine; shared hosting.
- Virtualization:
- Clouds

Servers typically share persistent storage, also in the cloud.

Cloud computing: pay by the number of instances that you've started up.

Providers offer different instance sizes: vary in cores, memory, GPU...



Need more computes? Launch an instance!

Input: Virtual Machine image.

Mechanics: use a command-line or web-based tool.

New instance gets an IP address and is network-accessible.
You have full root access to that instance.


Amazon provides public images:

- different Linux distributions;
- Windows Server; and
- OpenSolaris (maybe not anymore?).

You can build an image which contains software you want, including Hadoop and OpenMPI.

Presumably you don't want to pay forever for your instances.

When you're done with an instance:



All data on instance goes away.

To keep persistent results:

- mount a storage device, also on the cloud (e.g. Amazon Elastic Block Storage); or,
- connect to a database on a persistent server (e.g. Amazon SimpleDB or Relational Database Service); or,
- you can store files on the Web (e.g. Amazon S3).

Key idea: scaling to big data systems introduces substantial overhead.

Up next: Laptop vs. 128-core big data systems.

Are big data systems obviously good?
Have we measured (the right thing)?

The important metric is not just scalability;
absolute performance matters a lot.

Don't want: scaling up to n systems
to deal with complexity of scaling up to n .



Or, as Oscar Wilde put it:
“The bureaucracy is expanding to meet the needs of the expanding
bureaucracy.”

Compare: competent single-threaded implementation vs. top big data systems.

Domain: graph processing algorithms— PageRank and graph connectivity
(bottleneck is label propagation).

Subjects: graphs with billions of edges
(a few GB of data.)

Twenty pagerank iterations

System	cores	twitter_rv	uk_2007_05
Spark	128	857s	1759s
Giraph	128	596s	1235s
GraphLab	128	249s	833s
GraphX	128	419s	462s
Single thread	1	300s	651s

Label propagation to fixed-point (graph connectivity)

System	cores	twitter_rv	uk_2007_05
Spark	128	1784s	8000s+
Giraph	128	200s	8000s+
GraphLab	128	242s	714s
GraphX	128	251s	800s
Single thread	1	153s	417s

- “If you are going to use a big data system for yourself, see if it is faster than your laptop.”
- “If you are going to build a big data system for others, see that it is faster than my laptop.”

Movie Hour, featuring NoSQL Bane

Let's take a humorous look at cloud computing: James Mickens' session from Monitorama PDX 2014.



<https://vimeo.com/95066828>