

Lecture 6 — Working with Threads

Patrick Lam & Jeff Zarnett

`patrick.lam@uwaterloo.ca jzarnett@uwaterloo.ca`

Department of Electrical and Computer Engineering
University of Waterloo

December 5, 2018

- Available on most systems
- Windows has Pthreads Win32, but I wouldn't use it; use Linux for this course
- API available by `#include <pthread.h>`
- Compile with pthread flag
(`gcc -pthread prog.c -o prog`)

- `pthread_create`
- `pthread_exit`
- `pthread_join`
- `pthread_yield`
- `pthread_attr_init`
- `pthread_attr_destroy`
- `pthread_cancel`

- Now part of the C++ standard (library)
- API available with `#include <thread>`
- Compile with flags:
(`g++ -std=c++11 -pthread prog.c -o prog`)

```
int pthread_create(pthread_t* thread,
                  const pthread_attr_t* attr,
                  void* (*start_routine)(void*),
                  void* arg);
```

thread: creates a handle to a thread at pointer location

attr: thread attributes (NULL for defaults, more details later)

start_routine: function to start execution

arg: value to pass to start_routine

returns 0 on success, error number otherwise
(contents of *thread are undefined)

Creating Threads—Pthreads Example

```
#include <pthread.h>
#include <stdio.h>

void* run(void*) {
    printf("In_run\n");
}

int main() {
    pthread_t thread;
    pthread_create(&thread, NULL, run, NULL);
    printf("In_main\n");
}
```

Simply creates a thread and terminates
(usage isn't really right, as we'll see.)

Creating Threads—C++11 Example

```
#include <thread>
#include <iostream>

void run() {
    std::cout << "In_run\n";
}

int main() {
    std::thread t1(run);
    std::cout << "In_main\n";
    t1.join(); // hang in there...
}
```

```
int pthread_join(pthread_t thread,  
                 void** retval)
```

thread: wait for this thread to terminate (thread must be joinable).

retval: stores exit status of thread (set by `pthread_exit`) to the location pointed by `*retval`. If cancelled, returns `PTHREAD_CANCELED`. `NULL` is ignored.

returns 0 on success, error number otherwise.

Only call this one time per thread! Multiple calls on the same thread leads to undefined behaviour.

Waiting for Threads—Pthreads example

```
#include <pthread.h>
#include <stdio.h>

void* run(void*) {
    printf("In_run\n");
}

int main() {
    pthread_t thread;
    pthread_create(&thread, NULL, run, NULL);
    printf("In_main\n");
    pthread_join(thread, NULL);
}
```

This now waits for the newly created thread to terminate.

Creating Threads—C++11 Example

```
#include <thread>
#include <iostream>

void run() {
    std::cout << "In_run\n";
}

int main() {
    std::thread t1(run);
    std::cout << "In_main\n";
    t1.join(); // aha!
}
```

Passing Data to Pthreads threads...Wrongly

Consider this snippet:

```
int i;  
for (i = 0; i < 10; ++i)  
    pthread_create(&thread[i], NULL, run, (void*)&i);
```

This is a **terrible** idea. Why?

Passing Data to Pthreads threads...Wrongly

Consider this snippet:

```
int i;  
for (i = 0; i < 10; ++i)  
    pthread_create(&thread[i], NULL, run, (void*)&i);
```

This is a **terrible** idea. Why?

- 1 The value of `i` will probably change before the thread executes
- 2 The memory for `i` may be out of scope, and therefore invalid by the time the thread executes

Passing Data to Pthreads threads

Correct:

```
int i;  
int*  
for (i = 0; i < 10; ++i) {  
    arg = malloc( sizeof( int ) );  
    *arg = i;  
    pthread_create(&thread[i], NULL, run, arg);  
}
```

`int*` and `int` are always the same size, right guys?

What about:

```
int i;
for (i = 0; i < 10; ++i)
    pthread_create(&thread[i], NULL, run, (void*)i);

...

void* run(void* arg) {
    int id = (int)arg;
```

This is suggested in the book, but should carry a warning:

`int*` and `int` are always the same size, right guys?

What about:

```
int i;  
for (i = 0; i < 10; ++i)  
    pthread_create(&thread[i], NULL, run, (void*)i);  
  
...  
  
void* run(void* arg) {  
    int id = (int)arg;
```

This is suggested in the book, but should carry a warning:

- Beware size mismatches between arguments: no guarantee that a pointer is the same size as an `int`, so your data may overflow.
- Sizes of data types change between systems. For maximum portability, just use pointers you got from `malloc`.

It's easier to get data to threads in C++11:

```
#include <thread>
#include <iostream>

void run(int i) {
    std::cout << "In run_" << i << "\n";
}

int main() {
    for (int i = 0; i < 10; ++i) {
        std::thread t1(run, i);
        t1.detach(); // see the next slide...
    }
}
```

Getting Data from C++11 threads

...but it's harder to get data back.

Use async and future abstractions:

```
#include <thread>
#include <iostream>
#include <future>

int run() {
    return 42;
}

int main() {
    std::future<int> t1_retval =
        std::async(std::launch::async, run);
    std::cout << t1_retval.get();
}
```

Joinable threads (the default) wait for someone to call `pthread_join` before they release their resources.

Detached threads release their resources when they terminate, without being joined.

```
int pthread_detach(pthread_t thread);
```

thread: marks the thread as detached

returns 0 on success, error number otherwise.

Calling `pthread_detach` on an already detached thread results in undefined behaviour.

```
void pthread_exit(void *retval);
```

retval: return value passed to function that calls `pthread_join`

`start_routine` returning is equivalent to calling `pthread_exit` with that return value;

`pthread_exit` is called implicitly when the `start_routine` of a thread returns.

There is no C++11 equivalent.

By default, threads are *joinable* on Linux, but a more portable way to know what you're getting is to set thread attributes. You can change:

- Detached or joinable state
- Scheduling inheritance
- Scheduling policy
- Scheduling parameters
- Scheduling contention scope
- Stack size
- Stack address
- Stack guard (overflow) size

```
size_t stacksize;
pthread_attr_t attributes;
pthread_attr_init(&attributes);
pthread_attr_getstacksize(&attributes, &stacksize);
printf("Stack size = %i\n", stacksize);
pthread_attr_destroy(&attributes);
```

Running this on a laptop produces:

```
jon@riker examples master % ./ stack_size
Stack size = 8388608
```

Setting a thread state to joinable:

```
pthread_attr_setdetachstate(&attributes ,
                           PTHREAD_CREATE_JOINABLE);
```

Detached Threads: Warning!

```
#include <pthread.h>
#include <stdio.h>

void* run(void*) {
    printf("In_run\n");
}

int main() {
    pthread_t thread;
    pthread_create(&thread, NULL, run, NULL);
    pthread_detach(thread);
    printf("In_main\n");
}
```

When I run it, it just prints “In main”, why?

Detached Threads: Solution to Problem

```
#include <pthread.h>
#include <stdio.h>

void* run(void*) {
    printf("In_run\n");
}

int main() {
    pthread_t thread;
    pthread_create(&thread, NULL, run, NULL);
    pthread_detach(thread);
    printf("In_main\n");
    pthread_exit(NULL); // This waits for all detached
                       // threads to terminate
}
```

Make the final call `pthread_exit` if you have any detached threads. (There is no C++11 equivalent.)

- Be aware of scheduling (you can also set affinity with pthreads on Linux).
- Make sure the libraries you use are **thread-safe**:
 - Means that the library protects its shared data.
- glibc reentrant functions are also safe: a program can have more than one thread calling these functions concurrently.
- **Example:** rand_r versus rand.