Lecture 6 — Race Conditions & Synchronization

Patrick Lam & Jeff Zarnett

`p.lam@ece.uwaterloo.ca, jzarnett@uwaterloo.ca`

Department of Electrical and Computer Engineering
University of Waterloo

December 3, 2017

"Knock knock."
"Race Condition."
"Who's there?"

Define A Race Condition

A race occurs when you have two concurrent accesses to the same memory location, at least one of which is a **write**.

This definition is a little bit strict.

We could also say that there is a race condition if there is some form of output, such as writing to the console.

If one thread is going to write "1" to the console and another is going to write "2", then we could have a race condition.

If there is no co-ordination, we could get output of "12" or "21".

If the order here is unimportant, there's no issue; but if one order is correct, then the appearance of the other is a bug.

When there's a race, the final state may not be the same as running one access to completion and then the other.

But it "usually" is. It's nondeterministic.

The fact that the output is often "12" and only very occasionally "21" may make it very difficult to track down the source of the problem.

In other situations (e.g., processor design) these are sometimes referred to as data hazards or dependencies.

1 **RAW** (Read After Write)
2 **WAR** (Write After Read)
3 **WAW** (Write After Write)
4 **RAR** (Read After Read) - No such hazard!

It's a Kind of Hazard

Race conditions typically arise between variables shared between threads.

```c
#include <stdlib.h>
#include <stdio.h>
#include <pthread.h>

void* run1(void* arg) {
    int* x = (int*) arg;
    *x += 1;
}

void* run2(void* arg) {
    int* x = (int*) arg;
    *x += 2;
}

int main(int argc, char *argv[])
{
    int* x = malloc(sizeof(int));
    *x = 1;
    pthread_t t1, t2;
    pthread_create(&t1, NULL, &run1, x);
    pthread_join(t1, NULL);
    pthread_create(&t2, NULL, &run2, x);
    pthread_join(t2, NULL);
    printf("%d\n", *x);
    free(x);
    return EXIT_SUCCESS;
}
```

```
int main(int argc, char *argv[])
{
    int* x = malloc(sizeof(int));
    *x = 1;
    pthread_t t1, t2;
    pthread_create(&t1, NULL, &run1, x);
    pthread_create(&t2, NULL, &run2, x);
    pthread_join(t1, NULL);
    pthread_join(t2, NULL);
    printf("%d\n", *x);
    free(x);
    return EXIT_SUCCESS;
}
```

Now do we have a race?

What are the possible outputs? (Assume that initially ∗x is 1.) We'll look at compiler intermediate code (three-address code) to tell.

```
run1                    run2
D.1 = *x;               D.1 = *x;
D.2 = D.1 + 1;          D.2 = D.1 + 2
*x = D.2;               *x = D.2;
```

Memory reads and writes are key in data races.

Let's call the read and write from `run1` R1 and W1; R2 and W2 from `run2`.

Here are all possible orderings:

Order				*x
R1	W1	R2	W2	4
R1	R2	W1	W2	3
R1	R2	W2	W1	2
R2	W2	R1	W1	4
R2	R1	W2	W1	2
R2	R1	W1	W2	3

You'll need some sort of synchronization to get sane results from multithreaded programs.

We'll start by talking about how to use mutual exclusion in Pthreads.

Mutexes are the most basic type of synchronization. As a reminder:

- Only one thread can access code protected by a mutex at a time.
- All other threads must wait until the mutex is free before they can execute the protected code.

Here's an example of using mutexes:

PThreads	**C++11**

```
pthread_mutex_t m1_static = PTHREAD_MUTEX_INITIALIZER;    mutex m1;
pthread_mutex_t m2_dynamic;                               mutex * m2;

pthread_mutex_init(&m2_dynamic, NULL);                    m2 = new mutex();
...                                                       // ...
pthread_mutex_destroy(&m1_static);
pthread_mutex_destroy(&m2_dynamic);                       delete (m2);
```

You can initialize mutexes statically (as with `m1_static`) or dynamically (`m2_dynamic`).

If you want to include attributes, you need to use the dynamic version.

Both threads and mutexes use the notion of attributes.

- **Protocol**: specifies the protocol used to prevent priority inversions for a mutex.
- **Prioceiling**: specifies the priority ceiling of a mutex.
- **Process-shared**: specifies the process sharing of a mutex.

You can specify a mutex as *process shared* so that you can access it between processes.

PThreads

```
// code
pthread_mutex_lock(&m1);
// protected code
pthread_mutex_unlock(&m1);
// more code
```

C++11 Threads

```
// code
m1.lock();
// protected code
m1.unlock();
// more code
```

- Everything within the `lock` and `unlock` is protected.
- Be careful to avoid deadlocks if you are using multiple mutexes (always acquire locks in the same order across threads).
- Another useful primitive is `pthread_mutex_trylock`. later.

Why are we bothering with locks? Data races. A data race occurs when two concurrent actions access the same variable and at least one of them is a **write**. (This shows up on Assignment 1!)

```
static int counter = 0;

void* run(void* arg) {
    for (int i = 0; i < 100; ++i) {
        ++counter;
    }
}

int main(int argc, char *argv[]) {
    // Create 8 threads
    // Join 8 threads
    printf("counter = %i\n", counter);
}
```

```
static pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;
static int counter = 0;

void* run(void* arg) {
    for (int i = 0; i < 100; ++i) {
        pthread_mutex_lock(&mutex);
        ++counter;
        pthread_mutex_unlock(&mutex);
    }
}

int main(int argc, char *argv[]) {
    // Create 8 threads
    // Join 8 threads
    pthread_mutex_destroy(&mutex);
    printf("counter = %i\n", counter);
}
```

- Call `lock` on mutex ℓ_1. Upon return from `lock`, your thread has exclusive access to ℓ_1 until it `unlocks` it.
- Other calls to `lock` ℓ_1 will not return until `m1` is available.

Key idea: locks protect resources; only one thread can hold a lock at a time.

A second thread trying to obtain the lock (i.e. contending for the lock) has to wait, or block, until the first thread releases the lock.

So only one thread has access to the protected resource at a time.

The code between the lock acquisition and release is known as the critical region or critical section.

Some mutex implementations also provide a "try-lock" primitive.

This grabs the lock if it's available, or returns control to the thread if it's not.

This enables the thread to do something else. (Kind of like non-blocking I/O!)

Excessive use of locks can serialize programs.

Consider two resources A and B protected by a single lock ℓ.

Then a thread that's just interested in B still has acquire ℓ, which requires it to wait for any other thread working with A.

Example: Linux Big Kernel Lock

Spinlocks are a variant of mutexes, where the waiting thread repeatedly tries to acquire the lock instead of sleeping.

Use spinlocks when you expect critical sections to finish quickly.

Spinning for a long time consumes lots of CPU resources.

Many lock implementations use both sleeping and spinlocks: spin for a bit, then sleep longer.

When would we ever want to use a spinlock?

What we normally expect is to block until the lock becomes available.

But that means a process switch, and then a switch back in the future when the lock is available. This takes nonzero time.

It's optimal to use a spinlock if the amount of time we expect to wait for the lock is less than the amount of time it would take to do two process switches.

As long as we have a multicore system.

Two observations:

- If there are only reads, there's no datarace.
- Often, writes are relatively rare.

With mutexes/spinlocks, you have to lock the data, even for a read, since a write could happen.

But, most of the time, reads can happen in parallel, as long as there's no write.

Solution: Multiple threads can hold a read lock
    (`pthread_rwlock_rdlock`)
but only one thread may hold the associated write lock
    (`pthread_rwlock_wrlock`);
grabbing the write waits until current readers are done.

Semaphores have a `value`. You specify initial `value`.

Semaphores allow sharing of a # of instances of a resource.

Two fundamental operations: `wait` and `post`.

- `wait` is like `lock`; reserves the resource and decrements the value.
    - If value is 0, sleep until value is greater than 0.
- `post` is like `unlock`; releases the resource and increments the value.

Allows you to ensure that (some subset of) a collection of threads all reach the barrier before finishing.

Pthreads: A barrier is a `pthread_barrier_t`.

Functions: `_init()` (parameter: how many threads the barrier should wait for) and `_destroy()`.

Also `_wait()`: similar to `pthread_join()`, but waits for the specified number of threads to arrive at the barrier

We'll talk more about this in a few weeks.

Modern CPUs support atomic operations, such as compare-and-swap, which enable experts to write lock-free code.

Lock-free implementations are extremely complicated and must still contain certain synchronization constructs.

```
#include <semaphore.h>

int sem_init(sem_t *sem, int pshared, unsigned int value);
int sem_destroy(sem_t *sem);
int sem_post(sem_t *sem);
int sem_wait(sem_t *sem);
int sem_trywait(sem_t *sem);
```

- Also must link with `-pthread` (or `-lrt` on Solaris).
- All functions return `0` on success.
- Same usage as mutexes in terms of passing pointers.

How could you use as `semaphore` as a `mutex`?

```
#include <semaphore.h>

int sem_init(sem_t *sem, int pshared, unsigned int value);
int sem_destroy(sem_t *sem);
int sem_post(sem_t *sem);
int sem_wait(sem_t *sem);
int sem_trywait(sem_t *sem);
```

- Also must link with `-pthread` (or `-lrt` on Solaris).
- All functions return `0` on success.
- Same usage as mutexes in terms of passing pointers.

How could you use as `semaphore` as a `mutex`?

- If the initial `value` is 1 and you use `wait` to lock and `post` to unlock, it's equivalent to a `mutex`.

Here's an example from the book. How would you make this always print "Thread 1" then "Thread 2" using semaphores?

```
#include <pthread.h>
#include <stdio.h>
#include <semaphore.h>
#include <stdlib.h>

void* p1 (void* arg) { printf("Thread 1\n"); }

void* p2 (void* arg) { printf("Thread 2\n"); }

int main(int argc, char *argv[])
{
    pthread_t thread[2];
    pthread_create(&thread[0], NULL, p1, NULL);
    pthread_create(&thread[1], NULL, p2, NULL);
    pthread_join(thread[0], NULL);
    pthread_join(thread[1], NULL);
    return EXIT_SUCCESS;
}
```

Here's their solution. Is it actually correct?

```
sem_t sem;
void* p1 (void* arg) {
  printf("Thread 1\n");
  sem_post(&sem);
}
void* p2 (void* arg) {
  sem_wait(&sem);
  printf("Thread 2\n");
}

int main(int argc, char *argv[])
{
    pthread_t thread[2];
    sem_init(&sem, 0, /* value: */ 1);
    pthread_create(&thread[0], NULL, p1, NULL);
    pthread_create(&thread[1], NULL, p2, NULL);
    pthread_join(thread[0], NULL);
    pthread_join(thread[1], NULL);
    sem_destroy(&sem);
}
```

1. `value` is initially 1.

2. Say `p2` hits its `sem_wait` first and succeeds.

3. `value` is now 0 and `p2` prints "Thread 2" first.

- If `p1` happens first, it would just increase `value` to 2.

1. `value` is initially 1.

2. Say `p2` hits its `sem_wait` first and succeeds.

3. `value` is now 0 and `p2` prints "Thread 2" first.

- If `p1` happens first, it would just increase `value` to 2.

- Fix: set the initial `value` to 0.

  Then, if `p2` hits its `sem_wait` first, it will not print until `p1` posts (and prints "Thread 1") first.

- Used to notify the compiler that the variable may be changed by "external forces". For instance,

```
int i = 0;

while (i != 255) {
  ...
```

`volatile` prevents this from being optimized to:

```
int i = 0;

while (true) {
  ...
```

- Variable will not actually be `volatile` in the critical section and only prevents useful optimizations.
- Usually wrong unless there is a **very** good reason for it.