Tasks: OpenMP's thread-like mechanism.

The main new feature in OpenMP 3.0 is the notion of *tasks*. When the program executes a `#pragma omp task` statement, the code inside the task is split off as a task and scheduled to run sometime in the future. Tasks are more flexible than parallel sections, because parallel sections constrain exactly how many threads are supposed to run, and there is also always a join at the end of the parallel section. On the other hand, the OpenMP runtime can assign any task to any thread that's running. Tasks therefore have lower overhead.

<div align="center">

`#pragma omp` **task** *[clause [[,] clause]*]*

</div>

Generates a task for a thread in the team to run. When a thread enters the region it may:

- immediately execute the task; or

- defer its execution. (any other thread may be assigned the task)

Allowed Clauses: **if, final, untied, default, mergeable, private, firstprivate, shared**

`if` **and** `final` **Clauses.**

<div align="center">

if *(scalar-logical-expression)*

</div>

When expression is `false`, generates an undeferred task—
the generating task region is suspended until execution of the undeferred task finishes.

<div align="center">

final *(scalar-logical-expression)*

</div>

When expression is `true`, generates a final task.
All tasks within a final task are *included*.
Included tasks are undeferred and also execute immediately in the same thread.

Let's look at some examples of these clauses.

```
void foo () {
    int i;
    #pragma omp task if(0) // This task is undeferred
    {
        #pragma omp task
        // This task is a regular task
        for (i = 0; i < 3; i++) {
            #pragma omp task
            // This task is a regular task
```

```
            bar();
        }
    }
    #pragma omp task final(1) // This task is a regular task
    {
        #pragma omp task // This task is included
        for (i = 0; i < 3; i++) {
            #pragma omp task
            // This task is also included
            bar();
        }
    }
}
```

`untied` **and** `mergeable` **Clauses.**

untied

- A suspended task can be resumed by any thread.

- "untied" is ignored if used with **final**.

- Interacts poorly with thread-private variables and `gettid()`.

mergeable

- For an undeferred or included task, allows the implementation to generate a merged task instead.

- In a merged task, the implementation may re-use the environment from its generating task (as if there was no task directive).

For more: `docs.oracle.com/cd/E24457_01/html/E21996/gljyr.html`

```
#include <stdio.h>
void foo () {
    int x = 2;
    #pragma omp task mergeable
    {
        x++; // x is by default firstprivate
    }
    #pragma omp taskwait
    printf("%d\n",x); // prints 2 or 3
}
```

This is an incorrect usage of **mergeable**: the output depends on whether or not the task got merged. Merging tasks (when safe) produces more efficient code.

Taskyield.

#pragma omp **taskyield**

This directive specifies that the current task can be suspended in favour of another task.

Here's a good use of **taskyield**.

```
void foo (omp_lock_t * lock, int n) {
    int i;
    for ( i = 0; i < n; i++ )
    #pragma omp task
    {
        something_useful();
        while (!omp_test_lock(lock)) {
            #pragma omp taskyield
        }
        something_critical();
        omp_unset_lock(lock);
    }
}
```

Taskwait.

$$\text{\#pragma omp } \textbf{taskwait}$$

Waits for the completion of the current task's child tasks.

Two examples which show off tasks, from [ACD+09], include a web server (with unstructured requests) and a user interface which allows users to start tasks that are to run in parallel.

Here's pseudocode for the Boa webserver main loop from [ACD+09].

```
#pragma omp parallel
 /* a single thread manages the connections */
 #pragma omp single nowait
 while (!end) {
   process any signals
   foreach request from the blocked queue {
     if (request dependencies are met) {
       extract from the blocked queue
       /* create a task for the request */
       #pragma omp task untied
         serve_request(request);
     }
   }
   if (new connection) {
     accept_connection();
     /* create a task for the request */
     #pragma omp task untied
       serve_request(new connection);
   }
   select();
 }
```

The untied qualifier lifts restrictions on the task-to-thread mapping. All it means is that a task that's been started by one thread could be picked up by another and carried forward. This can have bad side effects if there's thread-private data involved. The single directive indicates that the runtime is only to use one thread to execute the next statement; otherwise, it could execute N copies of the statement, which does belong to a OpenMP parallel construct.

```
struct node {
    struct node *left;
    struct node *right;
};
extern void process(struct node *);

void traverse(struct node *p) {
```

3

```
    if (p->left) {
        #pragma omp task
        // p is firstprivate by default
        traverse(p->left);
    }
    if (p->right) {
        #pragma omp task
        // p is firstprivate by default
        traverse(p->right);
    }
    process(p);
}
```

If we want to guarantee a post-order traversal, we simply need to insert an explicit `#pragma omp taskwait` after the two calls to `traverse` and before the call to `process`.

Parallel Linked List Processing. We can spawn tasks to process linked list entries. It's hard to use two threads to traverse the list, though.

```
// node struct with data and pointer to next
extern void process(node* p);

void increment_list_items(node* head) {
    #pragma omp parallel
    {
        #pragma omp single
        {
            node * p = head;
            while (p) {
                #pragma omp task
                {
                    process(p);
                }
                p = p->next;
            }
        }
    }
}
```

Using Lots of Tasks. Let's see what happens if we spawn lots of tasks in a `single` directive.

```
#define LARGE_NUMBER 10000000
double item[LARGE_NUMBER];
extern void process(double);

int main() {
    #pragma omp parallel
    {
        #pragma omp single
        {
            int i;
            for (i=0; i<LARGE_NUMBER; i++) {
                #pragma omp task
                // i is firstprivate, item is shared
                process(item[i]);
            }
        }
    }
}
```

In this case, the main loop generates tasks, which are all assigned to the executing thread as it becomes available (because of `single`). When too many tasks get generated, OpenMP suspends the main thread, runs some tasks, then resumes the loop in the main thread.

Improved code.   It would be better to `untied` the spawned tasks, enabling them to run on multiple threads. Surround the for loop with #pragma omp task untied.

More Scoping Clauses

Besides the `shared`, `private` and `threadprivate`, OpenMP also supports `firstprivate` and `lastprivate`, which work like this.

Pthreads pseudocode for `firstprivate` clause:

```
int x;

void* run(void* arg) {
    int thread_x = x;
    // use thread_x
}
```

Pthread pseudocode for the `lastprivate` clause:

```
int x;

void* run(void* arg) {
    int thread_x;
    // use thread_x
    if (last_iteration) {
        x = thread_x;
    }
}
```

In other words, `lastprivate` makes sure that the variable x has the same value as if the loop executed sequentially.

`copyin` is like firstprivate, but for threadprivate variables.

Pthreads pseudocode for `copyin`:

```
int x;
int x[NUM_THREADS];

void* run(void* arg) {
  x[thread_num] = x;
  // use x[thread_num]
}
```

The `copyprivate` clause is only used with `single`. It copies the specified private variables from the thread to all other threads. It cannot be used with `nowait`.

Defaults.   default(shared) makes all variables shared; **default(none)** prevents sharing by default (creating compiler errors if you treat a variable as shared.)

Catch Me Outside...   A related problem with private variables is that sometimes you need access to them outside their parallel region. Here's some contrived code.

```
#include <omp.h>
#include <stdio.h>

int tid, a, b;

#pragma omp threadprivate(a)

int main(int argc, char *argv[])
{
    printf("Parallel_#1_Start\n");
    #pragma omp parallel private(b, tid)
    {
```

```
        tid = omp_get_thread_num();
        a = tid;
        b = tid;
        printf("T%d:_a=%d,_b=%d\n", tid, a, b);
    }

    printf("Sequential_code\n");
    printf("Parallel_#2_Start\n");
    #pragma omp parallel private(tid)
    {
        tid = omp_get_thread_num();
        printf("T%d:_a=%d,_b=%d\n", tid, a, b);
    }

    return 0;
}
```

This yields something like the following output:

```
% ./a.out
Parallel #1 Start
T6: a=6, b=6
T1: a=1, b=1
T0: a=0, b=0
T4: a=4, b=4
T2: a=2, b=2
T3: a=3, b=3
T5: a=5, b=5
T7: a=7, b=7
Sequential code
Parallel #2 Start
T0: a=0, b=0
T6: a=6, b=0
T1: a=1, b=0
T2: a=2, b=0
T5: a=5, b=0
T7: a=7, b=0
T3: a=3, b=0
T4: a=4, b=0
```

OpenMP Memory Model, Its Pitfalls, and How to Mitigate Them

OpenMP uses a **relaxed-consistency, shared-memory** model. This almost certainly doesn't do what you want. Here are its properties:

- All threads share a single store called *memory*—this store may not actually represent RAM.

- Each thread can have its own *temporary* view of memory.

- A thread's *temporary* view of memory is not required to be consistent with memory.

We'll talk more about memory models later. Now we're going to talk about the OpenMP model and why it's a problem.

Memory Model Pitfall. Consider this code.

```
                  a = b = 0
/* thread 1 */                    /* thread 2 */

atomic(b = 1) // [1]              atomic(a = 1) // [3]
atomic(tmp = a) // [2]            atomic(tmp = b) // [4]
if (tmp == 0) then               if (tmp == 0) then
    // protected section              // protected section
end if                           end if
```

6

Does this code actually prevent simultaneous execution? Let's reason about possible states.

Order				t1 tmp	t2 tmp
1	2	3	4	0	1
1	3	2	4	1	1
1	3	4	2	1	1
3	4	1	2	1	0
3	1	2	4	1	1
3	1	4	2	1	1

Looks like it (at least intuitively).

Sorry! With OpenMP's memory model, no guarantees: the update from one thread may not be seen by the other.

Restoring Sanity with Flush. We do rely on shared memory working "properly", but that's expensive. So OpenMP provides the **flush** directive.

<p style="text-align:center">#pragma omp flush [(list)]</p>

This directive makes the thread's temporary view of memory consistent with main memory; it:

- enforces an order on the memory operations of the variables.

The variables in the list are called the *flush-set*. If you give no variables, the compiler will determine them for you.

Enforcing an order on the memory operations means:

- All read/write operations on the *flush-set* which happen before the **flush** complete before the flush executes.
- All read/write operations on the *flush-set* which happen after the **flush** complete after the flush executes.
- Flushes with overlapping *flush-sets* can not be reordered.

To show a consistent value for a variable between two threads, OpenMP must run statements in this order:

1. t_1 writes the value to v;
2. t_1 flushes v;
3. t_2 flushes v also;
4. t_2 reads the consistent value from v.

Let's revise the example again.

```
               a = b = 0
/* thread 1 */                 /* thread 2 */

atomic(b = 1)                  atomic(a = 1)
flush(b)                       flush(a)
flush(a)                       flush(b)
atomic(tmp = a)                atomic(tmp = b)
if (tmp == 0) then             if (tmp == 0) then
    // protected section           // protected section
end if                         end if
```

OK. Will this now prevent simultaneous access?

Well, no.

The compiler can reorder the `flush(b)` in thread 1 or `flush(a)` in thread 2. If `flush(b)` gets reordered to after the protected section, we will not get our intended operation.

Correct Example.   We have to provide a list of variables to `flush` to prevent re-ordering:

```
                    a = b = 0
/* thread 1 */                    /* thread 2 */

atomic(b = 1)                     atomic(a = 1)
flush(a, b)                       flush(a, b)
atomic(tmp = a)                   atomic(tmp = b)
if (tmp == 0) then                if (tmp == 0) then
    // protected section              // protected section
end if                            end if
```

Where There Is Implicit Flush:

- `omp barrier`
- at entry to, and exit from, **omp critical**;
- at exit from **omp parallel**;
- at exit from **omp for**;
- at exit from **omp sections**;
- at exit from **omp single**.

Where There's No Implicit Flush:

- at entry to **for**;
- at entry to, or exit from, **master**;
- at entry to **sections**;
- at entry to **single**;
- at exit from **for**, **single** or **sections** with a **nowait**
    - **nowait** removes implicit flush along with the implicit barrier

This is not true for OpenMP versions before 2.5, so be careful.

Final thoughts on flush.   We've seen that it's very difficult to use flush properly. Really, you should be using mutexes or other synchronization instead of flush [Sue07], because you'll probably just get it wrong. But now you know what flush means.

Why Your Code is Slow

OpenMP code too slow? Avoid these pitfalls:

1. Unnecessary flush directives.

2. Using critical sections or locks instead of atomic.

3. Unnecessary concurrent-memory-writing protection:

    - No need to protect local thread variables.
    - No need to protect if only accessed in **single** or **master**.

4. Too much work in a critical section.

5. Too many entries into critical sections.

Example: Too Many Entries into Critical Sections.

```
#pragma omp parallel for
for (i = 0; i < N; ++i) {
    #pragma omp critical
    {
        if (arr[i] > max) max = arr[i];
    }
}
```

would be better as:

```
#pragma omp parallel for
for (i = 0 ; i < N; ++i) {
    #pragma omp flush(max)
    if (arr[i] > max) {
        #pragma omp critical
        {
            if (arr[i] > max) max = arr[i];
        }
    }
}
```

References

[ACD⁺09] Eduard Ayguadé, Nawal Copty, Alejandro Duran, Jay Hoeflinger, Yuan Lin, Federico Massaioli, Xavier Teruel, Priya Unnikrishnan, and Guansong Zhang. The design of OpenMP tasks. *IEEE Trans. Parallel Distrib. Syst.*, 20:404–418, March 2009. URL: http://portal.acm.org/citation.cfm?id=1512157.1512430.

[Sue07] Michael Suess. Please don't rely on memory barriers for synchronization, 2007. Online; accessed 12-December-2015. URL: http://www.thinkingparallel.com/2007/02/19/please-dont-rely-on-memory-barriers-for-synchronization/.