Lecture 14 — Memory Consistency

Patrick Lam & Jeff Zarnett
`p.lam@ece.uwaterloo.ca, jzarnett@uwaterloo.ca`

Department of Electrical and Computer Engineering
University of Waterloo

December 3, 2017

Sequential program: statements execute in order.
Your expectation for concurrency: sequential consistency.

> *"... the result of any execution is the same as if the operations of all the processors were executed in some sequential order, and the operations of each individual processor appear in this sequence in the order specified by its program." — Leslie Lamport*

In brief:

1 for each thread: in-order execution;
2 interleave the threads' executions.

No one has it: too expensive.
Recall the worked example for **flush** last time.

Memory Models: Sequential Consistency

Another view of sequential consistency:

- each thread induces an *execution trace*.
- always: program has executed some prefix of each thread's trace.

Compilers and processors may reorder non-interfering memory operations.

$$T1 : x = 1; r1 = y;$$

If two statements are independent:

- OK to execute them in either order.
- (equivalently: publish their results to other threads).

Reordering is a major compiler tactic to produce speedup.

Sequential consistency:

- No reordering of loads/stores.

Sequential consistency for datarace-free programs:

- If your program has no data races, then sequential consistency.

Relaxed consistency (only some types of reorderings):

- Loads can be reordered after loads/stores; and
- Stores can be reordered after loads/stores.

Weak consistency:

- Any reordering is possible.

Still, **reorderings** only allowed if they look safe in current context (i.e. independent; different memory addresses).

```
                      x = y = 0

/* thread 1 */                    /* thread 2 */
x = 1;                            y = x;
r1 = y;                           r2 = x;
```

Assume architecture not sequentially consistent
        (weak consistency).

Show me all possible (intermediate and final) memory values and how they
arise.

must include every permutation of lines
      (since they can be in any order);
then iterate over all the values.

Probably actually too long, but shows how memory reorderings complicate things.

When it can prove safety, the **compiler** may reorder instructions (not just the hardware).

Example: want thread 1 to print value set in thread 2.

```
                        f = 0

/* thread 1 */                      /* thread 2 */
while (f == 0) /* spin */;          x = 42;
printf("%d", x);                    f = 1;
```

- If thread 2 reorders its instructions, will we get our intended result?

    No.

When it can prove safety, the **compiler** may reorder instructions (not just the hardware).

Example: want thread 1 to print value set in thread 2.

```
                              f = 0

/* thread 1 */                      /* thread 2 */
while (f == 0) /* spin */;           x = 42;
printf("%d", x);                     f = 1;
```

- If thread 2 reorders its instructions, will we get our intended result?

  No.

A **memory fence** prevents memory operations from crossing the fence (also known as a **memory barrier**).

```
                         f = 0

/* thread 1 */                      /* thread 2 */
while (f == 0) /* spin */;          x = 42;
// memory fence                     // memory fence
printf("%d", x);                    f = 1;
```

- Now prevents reordering; get expected result.

Step 1: Don't use volatile on C/C++ variables [1].

Syntax depends on the compiler.

- Microsoft Visual Studio C++ Compiler:

```
_ReadWriteBarrier()
```

- Intel Compiler:

```
__memory_barrier()
```

- GNU Compiler:

```
__asm__ __volatile__ ("" ::: "memory");
```

The compiler also shouldn't reorder
across e.g. Pthreads mutex calls.

[1] http://stackoverflow.com/questions/78172/using-c-pthreads-do-shared-variables-need-to-be-volatile.

Just as an aside, here's gcc's inline assembly format

```
__asm__ ( assembler template
      : output operands                /* optional */
      : input operands                 /* optional */
      : list of clobbered registers    /* optional */
      );
```

Last slide used **__volatile__** with __asm__. This isn't the same as the normal C volatile. It means:

- The compiler may not reorder this assembly code
  and put it somewhere else in the program.

Memory Fences: Preventing Hardware Memory Reordering

Memory barrier: no access after the barrier becomes visible to the system (i.e. takes effect) until after all accesses before the barrier become visible.

Note: these are all x86 `asm` instructions.

`mfence`:

- All loads and stores before the fence finish before any more loads or stores execute.

`sfence`:

- All stores before the fence finish before any more stores execute.

`lfence`:

- All loads before the fence finish before any more loads execute.

Preventing Hardware Memory Reordering (Option 2)

Some compilers also support preventing hardware reordering:

- Microsoft Visual Studio C++ Compiler:

```
MemoryBarrier();
```

- Solaris Studio (Oracle) Compiler:

```
__machine_r_barrier();
__machine_w_barrier();
__machine_rw_barrier();
```

- GNU Compiler:

```
__sync_synchronize();
```

Fortunately, an OpenMP **flush** (or, better yet, mutexes) also preserve the order of variable accesses.

Stops reordering from both the compiler and hardware.

For GNU, flush is implemented as `__sync_synchronize();`

Note: proper use of memory fences makes `volatile` not very useful (again, `volatile` is not meant to help with threading, and will have a different behaviour for threading on different compilers/hardware).

Part I

Atomic Operations

We saw the **atomic** directive in OpenMP, plus C++11 atomics.

Most OpenMP atomic expressions map to
atomic hardware instructions.

Other atomic instructions exist.

Also called **compare and exchange** (`cmpxchg` instruction).

```
int compare_and_swap (int* reg, int oldval, int newval)
{
  int old_reg_val = *reg;
  if (old_reg_val == oldval)
     *reg = newval;
  return old_reg_val;
}
```

- Afterwards, you can check if it returned `oldval`.
- If it did, you know you changed it.

Use compare-and-swap to implement spinlock:

```
void spinlock_init(int* lock) { *lock = 0; }

void spinlock_lock(int* lock) {
    while(compare_and_swap(lock, 0, 1) != 0) {}
    __asm__ ("mfence");
}

void spinlock_unlock(int* lock) {
    __asm__ ("mfence");
    *lock = 0;
}
```

You'll see **cmpxchg** quite frequently in the Linux kernel code.

Sometimes you'll read a location twice.

If the value is the same, nothing has changed, right?

Sometimes you'll read a location twice.

If the value is the same, nothing has changed, right?

No. This is an **ABA problem**.

You can combat this by "tagging": modify value with nonce upon each write.

Can keep value separately from nonce; double compare and swap atomically swaps both value and nonce.

Memory ordering:

- Sequential consistency;
- Relaxed consistency;
- Weak consistency.

How to prevent memory reordering with fences.
Other atomic operations.

The ABA problem is not any sort of acronym nor a reference to this:
`https://www.youtube.com/watch?v=Sj_9CiNkkn4`

It's a value that is A, then changed to B, then changed back to A.

The ABA problem is a big mess for the designer of lock-free Compare-And-Swap routines.

1. P_1 reads A_i from location L_i.
2. P_k interrupts P_1; P_k stores the value B into L_i.
3. P_j stores the value A_i into L_i.
4. P_1 resumes; it executes a false positive CAS.

False Positive

It's a "false positive" because P_1's compare-and-swap operation succeeds even though the value at L_i has been modified in the meantime.

If this doesn't seem like a bad thing, consider this.

If you have a data structure that will be accessed by multiple threads, you might be controlling access to it by the compare-and-swap routine.

What should happen is the algorithm should keep trying until the data structure in question has not been modified by any other thread in the meantime.

But with a false positive we get the impression that things didn't change, even though they really did.

You can combat this by "tagging": modify value with nonce upon each write.

You can also keep the value separately from the nonce; double compare and swap atomically swaps both value and nonce.