Lecture 24 — Profiling

Patrick Lam
patrick.lam@uwaterloo.ca

Department of Electrical and Computer Engineering
University of Waterloo

December 9, 2019

Think back: what operations are fast and what operations are not?

Takeaway: our intuition is often wrong.

Not just at a macro level,
but at a micro level.

You may be able to narrow down that this computation of x is slow,
but if you examine it carefully... what parts of it are slow?

Programmers waste enormous amounts of time thinking about, or worrying about, the speed of noncritical parts of their programs, and these attempts at efficiency actually have a strong negative impact when debugging and maintenance are considered. We should forget about small efficiencies, say about 97% of the time: premature optimization is the root of all evil. Yet we should not pass up our opportunities in that critical 3%.

– Donald Knuth

Feeling lucky?
Maybe you optimized a slow part.

To make your programs or systems fast,
you need to find out what is currently slow
and improve it (duh!).

Up until now, it's mostly been about
        "let's speed this up".
We haven't taken much time to decide what
we should speed up.

General idea:

      collect data on what parts of the code
      are taking up most of the time.

- What functions get called?
- How long do functions take?
- What's using memory?

There is always the "informal" way:
"debug" by print statements.

On entry to `foo`,
log "entering function foo", plus timestamp.

On exit,
log "exiting foo", also with timestamp.

Kind of works, and [JZ] has used it to figure out what blocks of a single large function are taking a long time (updating exchange rates... yeah).

But this approach is not necessarily good.

This is "invasive" profiling—change the source code of the program to add instrumentation (log statements).

Plus we have to do a lot of manual accounting.

Doesn't really scale.

Like debugging:
if you get to be a wizard you can
maybe do it by code inspection.

But speculative execution inside your head
is harder for perf than debugging.

So: we want to use tools
and do this in a methodical way.

So far we've been looking at small problems.

Must **profile** to see what's slow in a large program.

Two main outputs:

- flat;
- call-graph.

Two main data gathering methods:

- statistical;
- instrumentation.

Flat Profiler:

- Only computes average time in a particular function.
- Does not include other (useful) information (callees).

Call-graph Profiler:

- Computes call times.
- Reports frequency of function calls.
- Gives a call graph: who called what function?

Statistical:

Mostly, take samples of the system state, that is:

- every 100ms, check the system state.
- will cause some slowdown, but not much.

Instrumentation:

Add instructions at specified program points:

- can do this at compile time or run time (expensive);
- can instrument either manually or automatically;
- like conditional breakpoints.

When writing large software projects:

- First, write clear and concise code.
  Don't do premature optimizations—
  focus on correctness.
- Profile to get a baseline of your performance:
  - allows you to easily track any performance changes;
  - allows you to re-design your program before it's too late.

Focus your optimization efforts on the code that matters.

Good signs:

- Time is spent in the right part of the system.
- Most time should not be spent handling errors; in non-critical code; or in exceptional cases.
- Time is not unnecessarily spent in OS.

Statistical profiler, plus some instrumentation for calls.

Runs completely in user-space.

Only requires a compiler.

Use the `-pg` flag with `clang`.

Run your program as you normally would.

- Your program will now create a `gmon.out` file.

Use gprof to interpret the results:
   gprof <executable>.

A program with 100 million calls to two math functions.

```
int main () {
    int i, x1=10, y1=3, r1=0;
    float x2=10, y2=3, r2=0;

    for(i=0;i<100000000;i++) {
        r1 += int_math(x1,y1);
        r2 += float_math(y2,y2);
    }
}
```

- Looking at the code, we have no idea what takes longer.
- Probably would guess floating point math taking longer.
- (Overall, silly example.)

```
int int_math(int x, int y){
    int r1;
    r1=int_power(x,y);
    r1=int_math_helper(x,y);
    return r1;
}

int int_math_helper(int x, int y){
    int r1;
    r1=x/y*int_power(y,x)/int_power(x,y);
    return r1;
}

int int_power(int x, int y){
    int i, r;
    r=x;
    for(i=1;i<y;i++){
        r=r*x;
    }
    return r;
}
```

```
float float_math (float x, float y) {
    float r1;
    r1=float_power (x,y);
    r1=float_math_helper (x,y);
    return r1;
}

float float_math_helper (float x, float y) {
    float r1;
    r1=x/y*float_power (y,x)/float_power (x,y);
    return r1;
}

float float_power (float x, float y){
    float i, r;
    r=x;
    for (i=1; i<y; i++) {
        r=r*x;
    }
    return r;
}
```

When we run the program, profile says:

```
Flat profile:

Each sample counts as 0.01 seconds.
  %    cumulative  self            self    total
 time   seconds   seconds   calls  ns/call ns/call  name
32.58     4.69      4.69  300000000   15.64   15.64  int_power
30.55     9.09      4.40  300000000   14.66   14.66  float_power
16.95    11.53      2.44  100000000   24.41   55.68  int_math_helper
11.43    13.18      1.65  100000000   16.46   45.78  float_math_helper
 4.05    13.76      0.58  100000000    5.84   77.16  int_math
 3.01    14.19      0.43  100000000    4.33   64.78  float_math
 2.10    14.50      0.30                            main
```

- One function per line.
- **% time:** the percent of the total execution time in this function.
- **self:** seconds in this function.
- **cumulative:** sum of this function's time + any above it in table.

```
Flat profile:

Each sample counts as 0.01 seconds.
  %    cumulative   self              self     total
 time    seconds   seconds    calls  ns/call  ns/call  name
32.58      4.69      4.69  300000000    15.64    15.64  int_power
30.55      9.09      4.40  300000000    14.66    14.66  float_power
16.95     11.53      2.44  100000000    24.41    55.68  int_math_helper
11.43     13.18      1.65  100000000    16.46    45.78  float_math_helper
 4.05     13.76      0.58  100000000     5.84    77.16  int_math
 3.01     14.19      0.43  100000000     4.33    64.78  float_math
 2.10     14.50      0.30                                main
```

- **calls:** number of times this function was called
- **self ns/call:** just self nanoseconds / calls
- **total ns/call:** average time for function execution, including any other calls the function makes

After the flat profile gives you a feel for which functions are costly, can get a better story from the call graph.

```
index % time   self  children    called       name
                                                <spontaneous>
[1]    100.0   0.30  14.19                     main [1]
               0.58   7.13 100000000/100000000     int_math [2]
               0.43   6.04 100000000/100000000     float_math [3]
---------------------------------------------------------------
               0.58   7.13 100000000/100000000     main [1]
[2]     53.2   0.58   7.13 100000000          int_math [2]
               2.44   3.13 100000000/100000000     int_math_helper [4]
               1.56   0.00 100000000/300000000     int_power [5]
---------------------------------------------------------------
               0.43   6.04 100000000/100000000     main [1]
[3]     44.7   0.43   6.04 100000000          float_math [3]
               1.65   2.93 100000000/100000000     float_math_helper [6]
               1.47   0.00 100000000/300000000     float_power [7]
```

The line with the index is the current function being looked at **(primary line)**.

- Lines above are functions which called this function.
- Lines below are functions which were called by this function (children).

Primary Line

- **time:** total percentage of time spent in this function and its children
- **self:** same as in flat profile
- **children:** time spent in all calls made by the function
  - should be equal to self + children of all functions below

Callers (functions above the primary line)

- **self:** time spent in primary function, when called from current function.
- **children:** time spent in primary function's children, when called from current function.
- **called:** number of times primary function was called from current function / number of nonrecursive calls to primary function.

Callees (functions below the primary line)

- **self:** time spent in current function when called from primary.
- **children:** time spent in current function's children calls when called from primary.
  - self + children is an estimate of time spent in current function when called from primary function.
- **called:** number of times current function was called from primary function / number of nonrecursive calls to current function.

index	% time	self	children	called	name
		2.44	3.13	100000000/100000000	int_math [2]
[4]	38.4	2.44	3.13	100000000	int_math_helper [4]
		3.13	0.00	200000000/300000000	int_power [5]
		1.56	0.00	100000000/300000000	int_math [2]
		3.13	0.00	200000000/300000000	int_math_helper [4]
[5]	32.4	4.69	0.00	300000000	int_power [5]
		1.65	2.93	100000000/100000000	float_math [3]
[6]	31.6	1.65	2.93	100000000	float_math_helper [6]
		2.93	0.00	200000000/300000000	float_power [7]
		1.47	0.00	100000000/300000000	float_math [3]
		2.93	0.00	200000000/300000000	float_math_helper [6]
[7]	30.3	4.40	0.00	300000000	float_power [7]

We can now see where most of the time comes from, and pinpoint locations that make unexpected calls, etc.

This example isn't too exciting; could simplify the math.

Part I

gperftools

Google Performance Tools include:

- CPU profiler.
- Heap profiler.
- Heap checker.
- Faster (multithreaded) `malloc`.

We'll mostly use the CPU profiler:

- purely statistical sampling;
- no recompilation; at most linking; and
- built-in visual output.

You can use the profiler without any recompilation.

- Not recommended—worse data.

```
LD_PRELOAD="/usr/lib/libprofiler.so" \
CPUPROFILE=test.prof ./test
```

The other option is to link to the profiler:

- -lprofiler

Both options read the CPUPROFILE environment variable:

- states the location to write the profile data.

You can use the profiling library directly as well:

- `#include <gperftools/profiler.h>`

Bracket code you want profiled with:

- `ProfilerStart()`
- `ProfilerStop()`

You can change the sampling frequency with the `CPUPROFILE_FREQUENCY` environment variable.

- **Default value:** 100

Like `gprof`, it will analyze profiling results.

```
% pprof test test.prof
    Enters "interactive" mode
% pprof --text test test.prof
    Outputs one line per procedure
% pprof --gv test test.prof
     Displays annotated call-graph via 'gv'
% pprof --gv --focus=Mutex test test.prof
    Restricts to code paths including a .*Mutex.* entry
% pprof --gv --focus=Mutex --ignore=string test test.prof
    Code paths including Mutex but not string
% pprof --list=getdir test test.prof
    (Per-line) annotated source listing for getdir()
% pprof --disasm=getdir test test.prof
    (Per-PC) annotated disassembly for getdir()
```

Can also output `dot`, `ps`, `pdf` or `gif` instead of `gv`.

Similar to the flat profile in `gprof`

```
jon@riker examples master % pprof ——text test test.prof
Using local file test.
Using local file test.prof.
Removing killpg from all stack traces.
Total: 300 samples
      95   31.7%   31.7%      102   34.0% int_power
      58   19.3%   51.0%       58   19.3% float_power
      51   17.0%   68.0%       96   32.0% float_math_helper
      50   16.7%   84.7%      137   45.7% int_math_helper
      18    6.0%   90.7%      131   43.7% float_math
      14    4.7%   95.3%      159   53.0% int_math
      14    4.7%  100.0%      300  100.0% main
       0    0.0%  100.0%      300  100.0% __libc_start_main
       0    0.0%  100.0%      300  100.0% _start
```

Columns, from left to right:

Number of checks (samples) in this function.
Percentage of checks in this function.

- Same as **time** in `gprof`.

Percentage of checks in the functions printed so far.

- Equivalent to **cumulative** (but in %).

Number of checks in this function and its callees.
Percentage of checks in this function and its callees.
Function name.

/tmp/profiler2_unittest
Total samples: 202
Focusing on: 202
Dropped nodes with <= 1 abs(samples)
Dropped edges with <= 0 samples

test_main_thread
155 (76.7%)
of 200 (99.0%)

Graphical Output Explained

Output was too small to read on the slide.

- Shows the same numbers as the text output.
- Directed edges denote function calls.
- Note: number of samples in callees = number in "this function + callees" - number in "this function".
- **Example:**
  `float_math_helper`, 51 (local) of 96 (cumulative)
  96 - 51 = 45 (callees)
    - callee `int_power` = 7 (bogus)
    - callee `float_power` = 38
    - callees total = 45

Call graph is not exact.

- In fact, it shows many bogus relations which clearly don't exist.
- For instance, we know that there are no cross-calls between `int` and `float` functions.

As with `gprof`, optimizations will change the graph.

You'll probably want to look at the text profile first, then use the `-focus` flag to look at individual functions.