Lecture 26 — Liar, Liar

Patrick Lam & Jeff Zarnett
patrick.lam@uwaterloo.ca jzarnett@uwaterloo.ca

Department of Electrical and Computer Engineering
University of Waterloo

December 23, 2019

Let's open with a video that illustrates one of the problems with sampling-based profiling:

`https://www.youtube.com/watch?v=jQDjJRYmeWg`

Is this fake?

Part I

Lies about Calling Context

Who can we trust?

Some profiler results are real.
Other results are interpolated, and perhaps wrong.

Reference: Yossi Kreinin,
`http://www.yosefk.com/blog/how-profilers-lie-the-cases-of-gprof-and-kcachegrind.html`

```
void work(int n) {
  volatile int i=0; //don't optimize away
  while(i++ < n);
}
void easy() { work(1000); }
void hard() { work(1000*1000*1000); }
int main() { easy(); hard(); }
```

```
[plam@lynch L27]$ gprof ./try gmon.out
Flat profile:

Each sample counts as 0.01 seconds.
  %   cumulative   self              self     total
 time   seconds   seconds    calls  ms/call  ms/call  name
101.30     1.68      1.68        2   840.78   840.78  work
  0.00     1.68      0.00        1     0.00   840.78  easy
  0.00     1.68      0.00        1     0.00   840.78  hard
```

That's not right!
easy takes $\approx$ 0s, hard takes 1.68s.

Need to understand how `gprof` works.

- **profil()**: asks glibc to record which instruction is currently executing ($100\times$/second).
- **mcount()**: records call graph edges; called by `-pg` instrumentation.

profil information is statistical;
mcount information is exact.

```
[plam@lynch L27]$ gprof ./try gmon.out
Flat profile:

Each sample counts as 0.01 seconds.
  %   cumulative   self              self     total
 time   seconds   seconds    calls  ms/call  ms/call  name
101.30     1.68     1.68        2   840.78   840.78  work
  0.00     1.68     0.00        1     0.00  *840.78* easy
  0.00     1.68     0.00        1     0.00  *840.78* hard
```

- calls: reliable;
- self seconds: sampled, but OK here;
- total ms/call: interpolated!

gprof sees:

- total of 1.68s in `work`,
- 1 call to `work` from `easy`;
- 1 call to `work` from `hard`.

All of these numbers are reliable.

gprof's unreliable conclusion:
`easy`, `hard` both cause 840ms of `work` time.

Wrong: `work` takes 1000000× longer when called from `hard`!

- contribution of children to parents;
- total runtime spent in self+children;
- etc.

Two cases:

- functions with only one caller
  (e.g. `f()` only called by `g()`); or,
- functions which always take the same time to complete
  (e.g. `rand()`).

What's sketchy:
Any function whose running time depends on its inputs,
and which is called from multiple contexts.

KCacheGrind is a frontend to callgrind.

callgrind is part of valgrind,
and runs the program under an x86 JIT.

Yes, hard takes all the time.

```
void worker1(int n) {
  volatile int i=0;
  while(i++<n);
}
void worker2(int n) {
  volatile int i=0;
  while(i++<n);
}
void manager(int n1, int n2) {
  worker1(n1);
  worker2(n2);
}
void project1() {
  manager(1000, 1000000);
}
void project2() {
  manager(1000000, 1000);
}
int main() {
  project1();
  project2();
}
```

Now `worker2` takes all the time in `project1`,
and `worker1` takes all the time in `project2`.
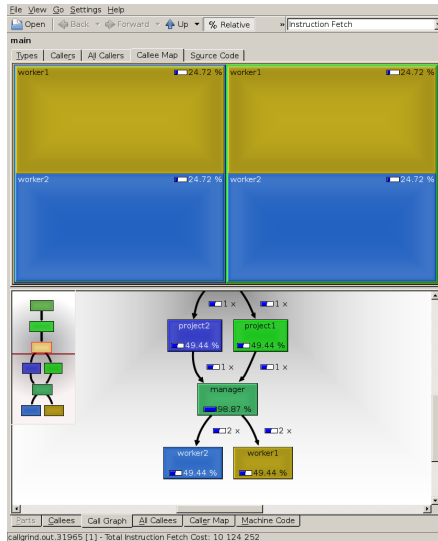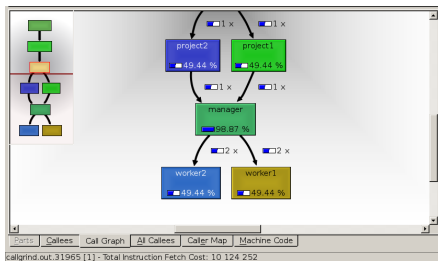
This is the call graph.
`worker1` and `worker2` do each take about 50% of time.
So do `project2` and `project1`.

(gprof would interpolate that too.)

KCacheGrind is reporting:

- `worker1` and `worker2` doing
  half the work in each project.

That's not what the code says.

- gprof reports time spent in `f()` and `g()`,
  and how many times `f()` calls `g()`.
- callgrind also reports
  time spent in `g()` when called from `f()`,
  i.e. some calling-context information.

- callgrind does *not* report time spent in `g()`
  when called from `f()` when called from `h()`.

  We don't get the `project1` to `manager` to `worker1` link.

- (We have Edges but need Edge-Pairs).

Some results are exact;
some results are sampled;
some results are interpolated.

If you understand the tool,
you understand where it can go wrong.

Understand your tools!

Part II

Lies from Metrics

While app-specific metrics can lie too,
mostly we'll talk about CPU perf counters.

Reference: Paul Khuong,
http://www.pvk.ca/Blog/2014/10/19/performance-optimisation-~-writing-an-essay/

We've talked about `mfence`.
Used in spinlocks, for instance.

Profiles said: spinlocking didn't take much time.
Empirically: eliminating spinlocks = better than expected!

Next step: create microbenchmarks.

Memory accesses to uncached locations,
or computations,

surrounded by store pairs/mfence/locks.

Use perf to evaluate impact of mfence vs lock.

```
$ perf annotate −s cache_misses
[...]
   0.06 :          4006b0:       and     %rdx,%r10
   0.00 :          4006b3:       add     $0x1,%r9
   ;; random (out of last level cache) read
   0.00 :          4006b7:       mov     (%rsi,%r10,8),%rbp
  30.37 :          4006bb:       mov     %rcx,%r10
   ;; foo is cached, to simulate our internal lock
   0.12 :          4006be:       mov     %r9,0x200fbb(%rip)
   0.00 :          4006c5:       shl     $0x17,%r10
   [... Skipping arithmetic with < 1% weight in the profile]
   ;; locked increment of an in−cache "lock" byte
   1.00 :          4006e7:       lock incb 0x200d92(%rip)
  21.57 :          4006ee:       add     $0x1,%rax
   [...]
   ;; random out of cache read
   0.00 :          400704:       xor     (%rsi,%r10,8),%rbp
  21.99 :          400708:       xor     %r9,%r8
   [...]
   ;; locked in−cache decrement
   0.00 :          400729:       lock decb 0x200d50(%rip)
  18.61 :          400730:       add     $0x1,%rax
   [...]
   0.92 :          400755:       jne     4006b0 <cache_misses+0x30>
```

Reads take 30 + 22 = 52% of runtime
Locks take 19 + 21 = 40%.

```
$ perf annotate -s cache_misses
[...]
    0.00 :        4006b0:        and    %rdx,%r10
    0.00 :        4006b3:        add    $0x1,%r9
    ;; random read
    0.00 :        4006b7:        mov    (%rsi,%r10,8),%rbp
   42.04 :        4006bb:        mov    %rcx,%r10
    ;; store to cached memory (lock word)
    0.00 :        4006be:        mov    %r9,0x200fbb(%rip)
    [...]
    0.20 :        4006e7:        mfence
    5.26 :        4006ea:        add    $0x1,%rax
    [...]
    ;; random read
    0.19 :        400700:        xor    (%rsi,%r10,8),%rbp
   43.13 :        400704:        xor    %r9,%r8
    [...]
    0.00 :        400725:        mfence
    4.96 :        400728:        add    $0x1,%rax
    0.92 :        40072c:        add    $0x1,%rax
    [...]
    0.36 :        40074d:        jne    4006b0 <cache_misses+0x30>
```

Looks like the reads take 85% of runtime,
while the mfence takes 15% of runtime.

Must also look at total # of cycles.

No atomic/fence: 2.81e9 cycles
lock inc/dec: 3.66e9 cycles
mfence: 19.60e9 cycles

That 15% number is a total lie.

- mfence underestimated;
- lock overestimated.

Why?

mfence = pipeline flush,
costs attributed to instructions being flushed.

Suppose we have a task that's going to get distributed over multiple computers (like a search).

If we look at the latency distribution, the problem is mostly that we see a long tail of events.

When we are doing a computation or search where we need all the results, we can only go as the slowest step.

1. Found in RAM
2. Disk Cache
3. Disk
4. and above... very strange!

Answer: CPU throttling!

This was happening on 25% of disk servers at Google, for an average of half an hour a day!

Another problem with sampling, this time from Lucene:

Why is perf limited to 100 KHz?

Answer: perf samples are done with interrupts (slow).

If you crank up the rate of interrupts, before long, you are spending all your time handling the interrupts rather than doing useful work.

SHIM gets around this by being more invasive.

This produces a bunch of data which can be dealt with later.