

ECE 459: Programming for Performance (Winter 2021)

Department of Electrical & Computer Engineering
University of Waterloo

About the Course

Undergraduate Calendar Description “Profiling computer systems; bottlenecks, Amdahl’s law. Concurrency: threads and locks. Techniques for programming multicore processors; cache consistency. Transactional memory. Streaming architectures, vectorization, and SIMD. High-performance programming languages.”

Prerequisites ECE 252, ECE 254, or SE 350 (or equivalent); Level at least 4A Electrical Engineering or Computer Engineering or Software Engineering.

Or, to put that in less formal terms: remember when we talked about semaphores and mutexes and all of that? If not, well, <https://github.com/jzarnett/ece252/tree/master/lectures>. (Or you can look at CS 343 material). Also, I really hope you feel comfortable with programming. This isn’t the place for you if you’re not feeling confident in programming. Really. We’re going to learn a new language (Rust!) so if you don’t feel comfortable with the fundamentals of programming and if learning a new language is frightening then flee while you can!

Brief Overview

Many modern software systems must process large amounts of data, either in the form of huge data sets or vast numbers of (concurrent) transactions. This course introduces students to techniques for profiling, rearchitecting, and implementing software systems that can handle industrial-sized inputs. These techniques will enable you to design and build critical software infrastructure, especially in an age of Big Data.

While you may have seen some of these ideas in the context of operating systems (ECE254/SE350/CS350) and concurrency (ECE252/CS343), this course gives you tools to make code run faster. The focus in OS/concurrency is understanding and implementing the primitives; our focus is on using them effectively.

We will sometimes see implementation details that you need to get right to write certain applications, but as with any university-level course, this course focusses more on the concepts than magic invocations, so that you can continue to apply the basic ideas after the technologies inevitably change.

General Information

Course Website The lecture material will be at <https://github.com/jzarnett/ece459>. The primary source for course materials is github. We’ll also use Piazza. See the section about the piazza policy.

Scheduling This course will be asynchronous. The lecture materials will be available in video format to consume at a time that's convenient for you. We still expect you to watch the lectures at a pace that is at least what would happen during in-person classes (so, about 2.5 hours per week).

The final assessment will happen during what would normally be the final exam period.

Schedule Oddities: Reading week is 17-21 February. In principle, 3rd and 4th year classes don't happen on the day of Capstone Project Symposia. There are a lot of programs represented in our course, to the point where no single program accounts for more than half of the students. Therefore we are not cancelling classes. We understand if you are going to miss it. We're also aware that IRS is a thing (8 February 2020).

Final Exam: The final exam dates are 8 April - 25 April. It could fall at any time in this period. Note that student travel plans are *not* considered an acceptable reason for missing an exam. When it is announced, please alert your instructor immediately if you have a conflict.

About Prof. Zarnett. I graduated from the Computer Engineering program at Waterloo (under a previous curriculum), and have since earned my Master's Degree (also at Waterloo) and my P.Eng. license. For the last 7+ years I have also been teaching here at UW and the other courses I have taught and worked on include ECE 150, ECE 155, ECE 252, ECE 254, ECE 290, ECE 356, and MTE 241.

In addition to being your instructor, I work full time in software engineering in industry. The best way to get questions answered is via Piazza. There's also the possibility of asking questions at office hours.

About the Teaching Assistants. Teaching assistants can help you with the course material, including assignments and exams. They will do most of the lab marking.

Teaching Assistants:

- Sakib Mohammad Chowdhury sm6chowd@uwaterloo.ca
- Aliasghar Iman aiman@uwaterloo.ca
- Michael Giannikouris mgiannik@uwaterloo.ca
- Stephen Li stephen.li@uwaterloo.ca

Textbook

There's no required textbook for the course. The notes cite a number of sources, some of which are useful books.

Evaluation

This course includes assignments, a midterm, and a final examination. Too many fourth-year courses have projects. This tends to make the end of term insane. We won't contribute to that problem too much.

Assignments	60% (4 at 15% each)
Midterm	10%
Final exam	30%

A grade of INC (Incomplete) is to be assigned if you do not attempt all assignments. If the assignments are not completed within 4 months of the end of the course to the satisfaction of the course instructor, the INC (Incomplete) grade will be converted to FTC (Failure to Complete). If the assignments are completed successfully

within the required timeframe, the overall grade will be calculated using the rules for late assignment submission. Furthermore, any submitted assignment must be an honest attempt—submitting an empty project or similar is not adequate.

Exams. Exams will be open-book, open-notes, calculators without communication capability, no communication devices. The University rules say if you miss the final exam, without an acceptable reason, your grade in the class will be DNW—Did Not Write. This is very undesirable. Show up for the final exam.

Tutorials. This class has tutorials scheduled. 4th year students basically never attend tutorials so there's no point in using them, right?

Assignments. Since this course has “programming” in the title, you will be expected to write code for these assignments. Here is a projected list of assignments for this course. We plan to have 4 assignments as below, although they are subject to change:

1. Manual parallelization for servers with pthreads and asynchronous I/O;
2. Using compiler-provided automatic parallelization and OpenMP;
3. GPU programming with OpenCL;
4. Profiling or load balancing.

Assignment hand-in will be done via `git` using the university provided `git.uwaterloo.ca` service.

Your assignment code will be checked for plagiarism using MOSS (Measure Of Software Similarity) as well as some manual checks. You may request to opt out of the automatic screening by sending a formal written letter to your instructor explaining why; a meeting will then follow to discuss the subject with the instructor.

You'll have at least two weeks to do each assignment. Trying to place the assignment due dates in the term is hard; We try to avoid FYDP deadlines, not place them on top of midterms, keep them vaguely in sync with the course material covered so far, and finally we're not allowed to make the due date for them during reading week or after the end of term. Oh, and because there are so many students from so many programs, it is effectively impossible to pick a date that works for everyone. Apologies in advance. That's why there are grace days (see below).

Lateness. Also known as “Grace Days”. You have 4 days of lateness to use on submissions throughout the term. Each day you hand in something late consumes one of the days of lateness. The fifth day of lateness causes your lowest assignment mark to be halved, while the sixth day causes both assignment marks to be halved. If you hand in something and you have more than 6 days of lateness, I'll start converting marks to 0 and dropping the associated late days. You don't get any credit for unused late days.

For example, you may hand in A1 one day late, A2 two days late, A3 1 one day late, and everything else on time. Or you can hand A2 four days late, if you hand in everything else on time. Finally, if you hand in A1 3 days late, A2 1 day late, A3 3 days late, and everything else on time, We will either give you a 0 for A1, leaving you with 4 late days, or give you a 0 for A2, leaving you with 5 late days and causing your mark for A1 to be halved. We'll choose the option which gives you more marks.

Grace days are tracked in Learn for transparency.

Group work. You may discuss assignments with others, but we expect each of you to do each assignment independently. Acceptable collaboration includes discussing ideas and structures with others, as well as helping others debug their code. If your code is too close in structure to someone else's code, you are going to have a problem. The best way to avoid such problems is by (1) not sending your code around; and (2) not writing down anything

beyond general notes (pseudocode) about other peoples' code. We will follow UW's Policy 71 if I discover any cases of plagiarism (and we have).

We want to emphasize that we take the issue of plagiarism very seriously, and so does the University of Waterloo. If you are uncertain about this subject, please seek some guidance. There are many resources available to you. You can check the university policies, talk to the course instructor, ECE/SE undergrad office, et cetera.

Or, let's sum this up in two short instructions:

1. Acknowledge the work of others.
2. If you are uncertain, ask!

Piazza Policy. Piazza is a great tool for collaboration and it allows rapid but asynchronous exchange of information. It almost goes without saying, but we ask you to be respectful and polite when communicating via this medium and to assume the best about others, as tone can be difficult to interpret on a discussion forum.

Before posting your question, please consider whether this question can be answered by looking at either the course notes, lab manual, or documentation available on the internet (e.g., man pages). To help others find answers later, (1) screenshot-only posts are discouraged (since the content does not show up in search results) and (2) please restrict a single thread to one question (or a closely-related set of questions).

For the most part, we encourage (but do not require) you to make your question(s) public so that other students may benefit from the answer. If something was non-obvious or unclear to you, it was likely to others as well. If your question includes detailed design information or code excerpts related to a deliverable (e.g., an assignment), then making it a private question is necessary.

Please also keep in mind that course staff may not be able to answer your question immediately as they have many other responsibilities throughout the term (and that the frequency of questions tends to spike before due dates or exams). We encourage you to think through the problem before posting, but acknowledge that just like "rubber duck debugging" sometimes the act of writing down the nature of the problem provides insight into solving it.

Re-marking. If you believe that your grade on the midterm is incorrect, you may ask that it be re-marked. To request that a question be re-marked, you will need to submit your request in writing via e-mail to the instructor.

When you submit your request, it should include the following: (1) Your name and student ID number; (2) a clear indication of which question or part of the deliverable is to be re-marked; and (3) an explanation of why you believe the grade assigned was incorrect.

Requests for re-marking may be submitted any time before the final exam. Be forewarned, when a deliverable is being re-marked, your grade could go up, it could stay the same, or it could go down. You will be notified of the outcome.

Extra Credit In this class, there will be no opportunities to earn extra credit. Make-up assignments or examinations will not be offered under any circumstances.

Attendance & Illness Personal opinion on attending classes: it is usually a good idea to attend lectures. That said, attendance is not taken and not graded.

Some advice Professor Gebotys gave long ago: If you are tired, go sleep at home. Sleeping in the lecture doesn't work; you will get poor quality of sleep and you won't learn the material while you're asleep, either.

If you feel ill, you should seek appropriate medical attention. If you miss an exam for health reasons, you need a verification of illness form (and it has to be rated "severe" on the form). Forms can be completed by the physicians

at Health Services or other healthcare providers in the area. If you anticipate missing a deliverable deadline or an examination for a non-medical reason, you should contact your instructor as soon as you are aware of the problem. Given sufficient notice, alternate arrangements may be possible. Alternate arrangements are rare and at instructor discretion.

Laptop and Device Policy The human visual system has evolved to perceive saber-toothed tigers in the savannah. Fortunately, tigers are rare in Waterloo, Ontario (Geese, on the other hand...). Unfortunately, your classmates are still human and hence their attention will be drawn to flashing lights (or Facebook, or movies, or video games) in their peripheral vision. We encourage everyone to be respectful of their classmates and to not distract them.

Wise use of computers and the Internet can be helpful for fully engaging in class. You might want to try out some syntax, or you might want to look up constructors, or you might want to verify your instructor's somewhat outrageous-sounding claim.

To support the benefits of the Internet while reducing distractions, we will adopt the following policy in this class. We are asking that the first 2 rows of class be text-oriented: if using a device, use a command prompt or text editor, maximized to the whole screen. Paper is always good, of course. Mac OS X and UNIX command prompts are probably your best bet; for those of you on Windows, you can use the Windows Subsystem for Linux. From the command prompt, you can use compilers and text-mode web browsers (w3m, lynx, links/elinks, etc....) tmux may also be helpful in managing multiple terminal sessions. Being proficient with the terminal is a highly-useful skill for a Software Engineer.

We acknowledge that lectures are not always engaging. Instead of distracting screen content, we recommend non-distracting ways of tuning out, like doodling on paper (while taking notes), or doing homework. (we also recommend passing notes to each other instead of talking).

If you need to sit towards the front of the class and use a GUI program, then please discuss with your instructor to register yourself as an exception. If you sign up for the exception list, we'll ask to you agree to not display games, videos, or social media on your screen (unless it is part of the class).

Enforcement is a sensitive issue, especially given the existence of exceptions. We are primarily asking each of you to respect the policy on your own. But, if you see someone with games, videos, or social media in the terminal zone, you can politely bring it up with them.

tl;dr: paper or text-oriented programs in first 2 rows of class.

University Policies

Academic Integrity In order to maintain a culture of academic integrity, members of the University of Waterloo community are expected to promote honesty, trust, fairness, respect and responsibility. Check www.uwaterloo.ca/academicintegrity/ for more information.

Grievance A student who believes that a decision affecting some aspect of their university life has been unfair or unreasonable may have grounds for initiating a grievance. Read Policy 70, Student Petitions and Grievances, Section 4, adm.uwaterloo.ca/infosec/Policies/policy70.htm
If in doubt, contact the department's administrative assistant, who will provide further assistance.

Discipline A student is expected to know what constitutes academic integrity (see above section) to avoid committing an academic offence, and to take responsibility for their actions. A student who is unsure whether an action constitutes an offence, or who needs help in learning how to avoid offences (e.g., plagiarism, cheating) or about "rules" for group work/collaboration should seek guidance from the course instructor, academic advisor, or the undergraduate Associate Dean. For information on categories of offences and types of penalties, students should refer to Policy 71, Student Discipline, www.adm.uwaterloo.ca/infosec/Policies/policy71.htm . For

typical penalties check Guidelines for the Assessment of Penalties, see
www.adm.uwaterloo.ca/infosec/guidelines/penaltyguidelines.htm.

Appeals A decision made or penalty imposed under Policy 70 (Student Petitions and Grievances) (other than a petition) or Policy 71 (Student Discipline) may be appealed if there is a ground. A student who believes he/she has a ground for an appeal should refer to Policy 72 (Student Appeals)
www.adm.uwaterloo.ca/infosec/Policies/policy72.htm.

Privacy Questions about the collection, use, and disclosure of personal information by the University, should be directed to the Freedom of Information and Privacy Coordinator, Secretariat, University of Waterloo, 200 University Avenue West, Waterloo, Ontario, Canada N2L 3G1. The email address of the Freedom of Information and Privacy Coordinator is fippa@uwaterloo.ca. See also University of Waterloo Policy 19: Access to and Release of Student Information; Information and Privacy.
<https://uwaterloo.ca/secretariat/policies-procedures-guidelines/policy-19>

Note for Students with Special Needs The AccessAbility Services (formerly known as OPD) located in Needles Hall, collaborates with all academic departments to arrange appropriate accommodations for students with disabilities without compromising the academic integrity of the curriculum. If you require academic accommodations to lessen the impact of your disability, please register with the AccessAbility Services office at the beginning of each academic term.